```
student@student:~$ echo $SHELL
/bin/bash
student@student:~$ nano ~/.bashrc
student@student:~$ source ~/.bashrc
student@student:~$ cd DSBDAL
student@student:~/DSBDAL$ echo $PATH
/home/student/DSBDAL/spark-3.5.1-bin-hadoop3/bin:/home/student/Documents/DSBDALab/spark-3.5.1-bin-hadoop3/bin:/home/student/DSBDAL/spark-3.5.1-
bin-hadoop3/bin:/home/student/Documents/DSBDALab/spark-3.5.1-bin-hadoop3/bin:/home/student/.local/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:
/usr/bin:/sbin:/bin:/usr/games:/usr/local/games:/snap/bin:/snap/bin:/home/student/bin:/usr/lib/jvm/java-8-openjdk-amd64/bin:/opt/spark/bin:/opt
/spark/sbin:/opt/spark/bin:/opt/spark/sbin
student@student:~/DSBDAL$ spark-shell<WebLog_Processing.scala
26/04/10 11:05:40 WARN Utils: Your hostname, student resolves to a loopback address: 127.0.1.1; using 10.11.5.204 instead (on interface eno1)
26/04/10 11:05:40 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
26/04/10 11:05:49 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Spark context Web UI available at http://10.11.5.204:4040
Spark context available as 'sc' (master = local[*], app id = local-1775799354795).
Spark session available as 'spark'.
Welcome to
      ____              __
     / __/__  ___ _____/ /__
    _\ \/ _ \/ _ `/ __/  '_/
   /___/ .__/\_,_/_/ /_/\_\   version 3.2.0
      /_/

Using Scala version 2.12.15 (OpenJDK 64-Bit Server VM, Java 1.8.0_482)
Type in expressions to have them evaluated.
Type :help for more information.

scala> import org.apache.log4j.{Level, Logger}
import org.apache.log4j.{Level, Logger}

scala> import org.apache.spark.sql.{Column, SparkSession}
import org.apache.spark.sql.{Column, SparkSession}

scala> import org.apache.spark.sql.functions.{regexp_extract,sum,col,to_date,udf,to_timestamp,desc,dayofyear,year}
```

```
scala> import org.apache.spark.sql.functions.{regexp_extract,sum,col,to_date,udf,to_timestamp,desc,dayofyear,year}
import org.apache.spark.sql.functions.{regexp_extract, sum, col, to_date, udf, to_timestamp, desc, dayofyear, year}

scala>

scala> val spark = SparkSession.builder().appName("WebLog").master("local[*]").getOrCreate()
26/04/10 11:06:02 WARN SparkSession$Builder: Using an existing SparkSession; some spark core configurations may not take effect.
spark: org.apache.spark.sql.SparkSession = org.apache.spark.sql.SparkSession@155d00b5

scala> val base_df = spark.read.text("/home/student/DSBDAL/weblog.csv")
base_df: org.apache.spark.sql.DataFrame = [value: string]

scala> base_df.printSchema()
root
 |-- value: string (nullable = true)


scala>

scala> import spark.implicits._
import spark.implicits._

scala> //this will produce a dataframe with a single column called value

scala> val base_df = spark.read.text("/home/student/DSBDAL/weblog.csv")
base_df: org.apache.spark.sql.DataFrame = [value: string]

scala> base_df.printSchema()
root
 |-- value: string (nullable = true)


scala> //let's look at some of the data

scala> base_df.show(3,false)
+-------------------------------------------------------------+
|value                                                        |
```

```
scala> //let's look at some of the data

scala> base_df.show(3,false)
+-----------------------------------------------------------------+
|value                                                            |
+-----------------------------------------------------------------+
|IP,Time,URL,Staus                                                |
|10.128.2.1,[29/Nov/2017:06:58:55,GET /login.php HTTP/1.1,200   |
|10.128.2.1,[29/Nov/2017:06:59:02,POST /process.php HTTP/1.1,302|
+-----------------------------------------------------------------+
only showing top 3 rows


scala>

scala> /*
     |       Parsing the log file
     |    */
     | val parsed_df = base_df.select(regexp_extract($"value","""^([^(\s|,)]+)""",1).alias("host"),
     |     regexp_extract($"value","""^.*\[(\d\d/\w{3}/\d{4}:\d{2}:\d{2}:\d{2})""",1).as("timestamp"),
     |     regexp_extract($"value","""^.*\w+\s+([^\s]+)\s+HTTP.*""",1).as("path"),
     |     regexp_extract($"value","""^.*,([^\s]+)$""",1).cast("int").alias("status"))
parsed_df: org.apache.spark.sql.DataFrame = [host: string, timestamp: string ... 2 more fields]

scala> parsed_df.show(5,false)
+----------+--------------------+-----------------------+------+
|host      |timestamp           |path                   |status|
+----------+--------------------+-----------------------+------+
|IP        |                    |                       |null  |
|10.128.2.1|29/Nov/2017:06:58:55|/login.php             |200   |
|10.128.2.1|29/Nov/2017:06:59:02|/process.php           |302   |
|10.128.2.1|29/Nov/2017:06:59:03|/home.php              |200   |
|10.131.2.1|29/Nov/2017:06:59:04|/js/vendor/moment.min.js|200   |
+----------+--------------------+-----------------------+------+
only showing top 5 rows
```

```
|10.128.2.1|29/Nov/2017:06:59:03|/home.php              |200   |
|10.131.2.1|29/Nov/2017:06:59:04|/js/vendor/moment.min.js|200   |
+----------+--------------------+-----------------------+------+
only showing top 5 rows


scala> parsed_df.printSchema()
root
 |-- host: string (nullable = true)
 |-- timestamp: string (nullable = true)
 |-- path: string (nullable = true)
 |-- status: integer (nullable = true)


scala>

scala> /*
     |      Data cleaning
     |    */
     |   // check if the initial dataset contain any null values
     |   println("Number of bad row in the initial dataset : " + base_df.filter($"value".isNull).count())
Number of bad row in the initial dataset : 0

scala>

scala>   // lets see our parsed dataset

scala>   val bad_rows_df = parsed_df.filter($"host".isNull || $"timestamp".isNull || $"path".isNull || $"status".isNull)
bad_rows_df: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [host: string, timestamp: string ... 2 more fields]

scala>   println("Number of bad rows : " + bad_rows_df.count())
Number of bad rows : 219

scala>   // same result as the previous statement but with different syntax

scala>   //val bad_rows_df = parsed_df.filter($"host".isNull.or($"timestamp".isNull).or($"path".isNull)

scala>   // .or($"status".isNull)).count
```

```
scala>   // same result as the previous statement but with different syntax

scala>   //val bad_rows_df = parsed_df.filter($"host".isNull.or($"timestamp".isNull).or($"path".isNull)

scala>   // .or($"status".isNull)).count

scala>

scala>   // lets count number of null values in each column

scala>   // we create a function that convert null value to one and then we count the number of one value

scala>   def count_null(col_name: Column): Column = sum(col_name.isNull.cast("int")).alias(col_name.toString())
count_null: (col_name: org.apache.spark.sql.Column)org.apache.spark.sql.Column

scala>   val t = parsed_df.columns.map(col_name => count_null(col(col_name)))
t: Array[org.apache.spark.sql.Column] = Array(sum(CAST((host IS NULL) AS INT)) AS host, sum(CAST((timestamp IS NULL) AS INT)) AS timestamp, sum
(CAST((path IS NULL) AS INT)) AS path, sum(CAST((status IS NULL) AS INT)) AS status)

scala>   parsed_df.select(t: _*).show()
+----+---------+----+------+
|host|timestamp|path|status|
+----+---------+----+------+
|   0|        0|   0|   219|
+----+---------+----+------+


scala>

scala>   // So all the null values are in status column, let's check what does it contain

scala>   val bad_status_df = base_df.select(regexp_extract($"value","""([^\d]+)$""",1).as("bad_status")).filter($"bad_status".notEqual(""))
bad_status_df: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [bad_status: string]

scala>   println("Number of bad rows : " + bad_status_df.count())
Number of bad rows : 219
```

```
scala>   println("Number of bad rows : " + bad_status_df.count())
Number of bad rows : 219

scala>   // So the bad content correspond to error result, in our case this is just polluting our logs and our results

scala>   bad_status_df.show(5)
+--------------------+
|          bad_status|
+--------------------+
|   IP,Time,URL,Staus|
|chmod:,cannot,'a....|
|chmod:,cannot,'er...|
|rm:,cannot,'*.o':,No|
|rm:,cannot,'a.out...|
+--------------------+
only showing top 5 rows


scala>
     |   /*
     |        Fix the rows with null status
     |    */
     |
     |   // we have two option, replace the null value by some other meaningful value, or delete the whole line
     |   // we will go with the other option since those lines are with no value for us
     |   // we will use the na subpackage on a dataframe
     |   val cleaned_df = parsed_df.na.drop()
cleaned_df: org.apache.spark.sql.DataFrame = [host: string, timestamp: string ... 2 more fields]

scala>

scala>   // let's check that we don't have any null value

scala>   println("The count of null value : " + cleaned_df.filter($"host".isNull || $"timestamp".isNull || $"path".isNull|| $"status".isNull).c
ount())
The count of null value : 0

scala>   // count before and after
```

```
scala>   // count before and after

scala>   println("Before : " + parsed_df.count() + " | After : " + cleaned_df.count())
Before : 16008 | After : 15789

scala>
     |   /*
     |         Parsing the timestamp
     |    */
     |   // let's try to cast the timestamp column to date
     |   // surprised ! we got null value, that's because when spark is not able to convert a date value
     |   // it just return null
     |   cleaned_df.select(to_date($"timestamp")).show(2)
+------------------+
|to_date(timestamp)|
+------------------+
|              null|
|              null|
+------------------+
only showing top 2 rows


scala>

scala>   // Let's fix this by converting the timestamp column to the format spark knows

scala>   val month_map = Map("Jan" -> 1, "Feb" -> 2, "Mar" -> 3, "Apr" -> 4, "May" -> 5, "Jun" -> 6, "Jul" -> 7, "Aug" -> 8
     |     , "Sep" -> 9, "Oct" -> 10, "Nov" -> 11, "Dec" -> 12)
month_map: scala.collection.immutable.Map[String,Int] = Map(Nov -> 11, Jul -> 7, Mar -> 3, Jan -> 1, Oct -> 10, Dec -> 12, Feb -> 2, May -> 5,
Apr -> 4, Aug -> 8, Sep -> 9, Jun -> 6)

scala>   def parse_clf_time(s: String) ={
     |     "%3$s-%2$s-%1$s %4$s:%5$s:%6$s".format(s.substring(0,2),month_map(s.substring(3,6)),s.substring(7,11)
     |       ,s.substring(12,14),s.substring(15,17),s.substring(18))
     |   }
parse_clf_time: (s: String)String
```

```
ome(class[value[0]: string])),Some(class[value[0]: string]),None,true,true)

scala>   val logs_df = cleaned_df.select($"*",to_timestamp(toTimestamp($"timestamp")).alias("time")).drop("timestamp")
logs_df: org.apache.spark.sql.DataFrame = [host: string, path: string ... 2 more fields]

scala>   logs_df.printSchema()
root
 |-- host: string (nullable = true)
 |-- path: string (nullable = true)
 |-- status: integer (nullable = true)
 |-- time: timestamp (nullable = true)


scala>   logs_df.show(2)
+----------+------------+------+-------------------+
|      host|        path|status|               time|
+----------+------------+------+-------------------+
|10.128.2.1|  /login.php|   200|2017-11-29 06:58:55|
|10.128.2.1|/process.php|   302|2017-11-29 06:59:02|
+----------+------------+------+-------------------+
only showing top 2 rows


scala>   // We cache the dataset so the next action would be faster

scala>   logs_df.cache()
res29: logs_df.type = [host: string, path: string ... 2 more fields]

scala>
     |   //       ====<  Analysis walk-trough  >====

scala>

scala>   /*
     |         status column statistics
     |    */
     |   logs_df.describe("status").show()
+-------+------------------+
```

```
     |   logs_df.describe("status").show()
+-------+------------------+
|summary|            status|
+-------+------------------+
|  count|             15789|
|   mean|230.19469250744189|
| stddev| 50.05853522906924|
|    min|               200|
|    max|               404|
+-------+------------------+


scala>   /*
     |        HTTP status analysis
     |    */
     |   logs_df.groupBy("status").count().sort("status").show()
+------+-----+
|status|count|
+------+-----+
|   200|11330|
|   206|   52|
|   302| 3498|
|   304|  658|
|   404|  251|
+------+-----+


scala>

scala>   /*
     |        Frequent Hosts
     |    */
     |   logs_df.groupBy("host").count().filter($"count" > 10).show()
+----------+-----+
|      host|count|
+----------+-----+
|10.131.2.1| 1626|
|10.128.2.1| 4257|
```

```
|10.131.2.1| 1626|
|10.128.2.1| 4257|
|10.130.2.1| 4056|
|10.131.0.1| 4198|
|10.129.2.1| 1652|
+----------+-----+


scala>   /*
     |        Visualizing Paths
     |    */
     |   logs_df.groupBy("path").count().sort(desc("count")).show()
+--------------------+-----+
|                path|count|
+--------------------+-----+
|          /login.php| 3298|
|           /home.php| 2653|
|/js/vendor/modern...| 1417|
|                   /|  862|
|/contestproblem.p...|  467|
|  /css/normalize.css|  408|
|/css/bootstrap.mi...|  404|
|/css/font-awesome...|  399|
|      /css/style.css|  395|
|       /css/main.css|  394|
|/js/vendor/jquery...|  387|
|/bootstrap-3.3.7/...|  382|
|        /process.php|  317|
|        /contest.php|  249|
|        /archive.php|  246|
|/fonts/fontawesom...|  245|
|         /robots.txt|  224|
|       /img/ruet.png|  213|
|/bootstrap-3.3.7/...|  191|
|/js/vendor/moment...|  173|
+--------------------+-----+
only showing top 20 rows
```

```
only showing top 20 rows


scala>   /*
     |        Top Paths
     |    */
     |   logs_df.groupBy("path").count().sort(desc("count")).show(10)
+--------------------+-----+
|                path|count|
+--------------------+-----+
|          /login.php| 3298|
|           /home.php| 2653|
|/js/vendor/modern...| 1417|
|                   /|  862|
|/contestproblem.p...|  467|
|  /css/normalize.css|  408|
|/css/bootstrap.mi...|  404|
|/css/font-awesome...|  399|
|      /css/style.css|  395|
|       /css/main.css|  394|
+--------------------+-----+
only showing top 10 rows


scala>

scala>   //       ====< Analyzing Web Server Log File >=====

scala>

scala>   /*
     |        Top Ten Error Paths
     |    */
     |   logs_df.filter($"status" =!= 200).groupBy("path").count().sort(desc("count")).show(10)
+--------------------+-----+
|                path|count|
+--------------------+-----+
|           /home.php| 2167|
|                   /|  741|
|        /process.php|  317|
|         /robots.txt|  224|
|         /action.php|   83|
|/contestproblem.p...|   74|
|/js/vendor/jquery...|   73|
|/css/bootstrap.mi...|   72|
|/js/vendor/modern...|   72|
|       /css/main.css|   68|
+--------------------+-----+
only showing top 10 rows


scala>   /*
     |        Number of unique Hosts
     |    */
     |   val unique_host_count = logs_df.select("host").distinct().count()
unique_host_count: Long = 5

scala>   println("Unique hosts : %d".format(unique_host_count))
Unique hosts : 5

scala>   /*
     |        Number of Unique Daily Hosts :
     |    */
     |   val daily_hosts_df = logs_df.withColumn("day",dayofyear($"time")).withColumn("year",year($"time")).select("host","day","year").distinc
t().groupBy("day","year").count().sort("year","day").cache()
daily_hosts_df: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [day: int, year: int ... 1 more field]

scala>   daily_hosts_df.show(5)
+---+----+-----+
|day|year|count|
```

```
+---+----+-----+
|day|year|count|
+---+----+-----+
|311|2017|    1|
|312|2017|    5|
|313|2017|    5|
|314|2017|    5|
|315|2017|    5|
+---+----+-----+
only showing top 5 rows


scala>   /*
     |        Average Number of Daily Requests per Host
     |    */
     |   val total_req_per_day_df = logs_df.withColumn("day", dayofyear($"time")).withColumn("year", year($"time")).groupBy("day", "year").coun
t()
total_req_per_day_df: org.apache.spark.sql.DataFrame = [day: int, year: int ... 1 more field]

scala>   val avg_daily_request_per_host_df = total_req_per_day_df.join(daily_hosts_df,total_req_per_day_df("day") === daily_hosts_df("day")&& t
otal_req_per_day_df("year") === daily_hosts_df("year")).select(daily_hosts_df("day"),daily_hosts_df("year"),(total_req_per_day_df("count") /dai
ly_hosts_df("count")).alias("avg_req_per_host_per_day")).cache()
avg_daily_request_per_host_df: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [day: int, year: int ... 1 more field]

scala>   avg_daily_request_per_host_df.show(5)
+---+----+------------------------+
|day|year|avg_req_per_host_per_day|
+---+----+------------------------+
|335|2017|                    93.6|
|327|2017|                    76.0|
| 60|2018|      10.333333333333334|
|350|2017|      51.666666666666664|
| 46|2018|       6.666666666666667|
+---+----+------------------------+
only showing top 5 rows


scala>
```